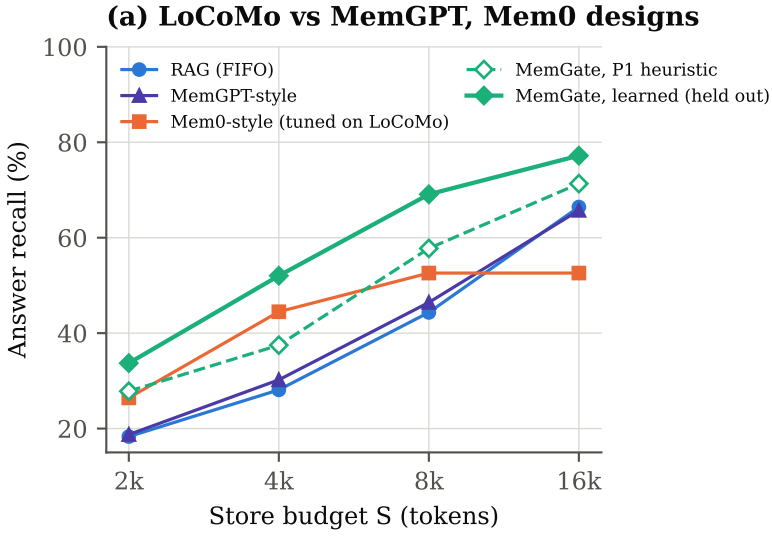

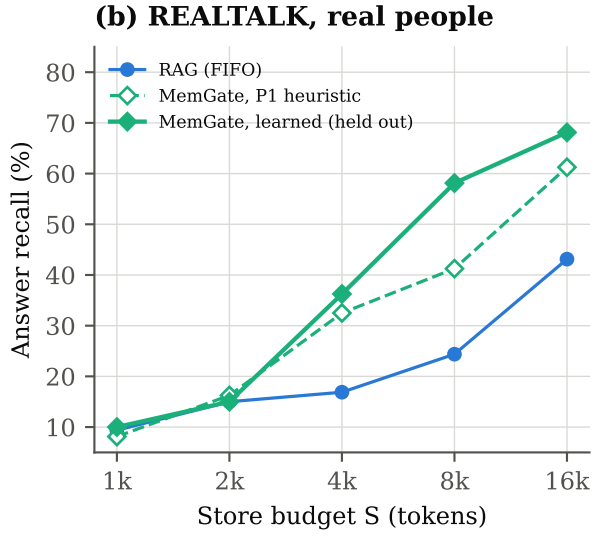

Context budget 2048. The learned scorer is fitted per held-out conversation (the other conversations plus the other corpus). The baselines' knobs were grid-searched on LoCoMo at S=4k, which favours them. LLM-free rebuilds of each system's memory management (memgate/external_baselines.py), not the products.